

RED-BLACK TREES

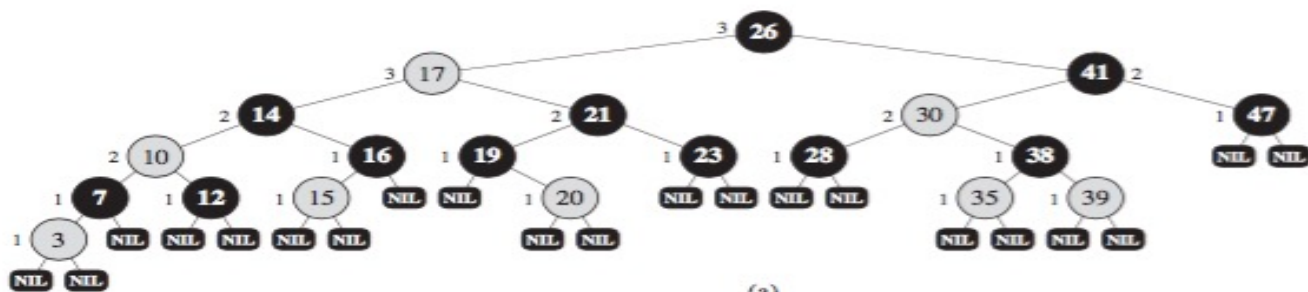
CS340

Red-Black Trees


- Red Black Trees are a kind of Binary Search Tree that are **balanced**
 - guarantee of logarithmic height $O(\lg n)$
 - Non-modifying binary-search-tree operations MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in $O(\lg n)$ time on red-black trees.
 - INSERT and DELETE are tricky! Must maintain red-black tree properties.
- R-B Trees have an extra bit of information at each node
 - the color: **Red?** or Black?

Red-Black Tree Properties


1. Every node is either **red** or black.
2. The root is black.
3. Every leaf is NIL and black.
4. If a node is **red** then both its children are black.
5. For each node, all simple paths to its descendant leaves have the same number of black nodes.
 - This number defines the black height of a node x , $bh(x)$.
 - By the previous property, $bh(x) \geq h(x)/2$



(a)




(b)



(c)

Height and Black Height

- Height: number of edges in longest path to a leaf
- Black Height: number of black nodes (including T.nil) on the path from x to leaf, not counting x (see property 5)



Red-black tree properties

Any node with height h has black-height, $bh(x) \geq h/2$.

Proof


Property 4: If a node is red then both its children are black.

This means that each red node must be followed by a black node. The reverse is not necessarily true.

Therefore, $\leq h/2$ nodes on the path from the node to a leaf are red. Hence $\geq h/2$ nodes are black.


Interview Questions

- Draw the red-black tree that results after TREE-INSERT is called on the tree below with key 36. If the inserted node is colored red, is the resulting tree a red-black tree? What if it is colored black?




THE BIG PICTURE

Rotations of trees




- Rotation to change shape
- One side gets taller, the other shrinks



Rotation of Trees

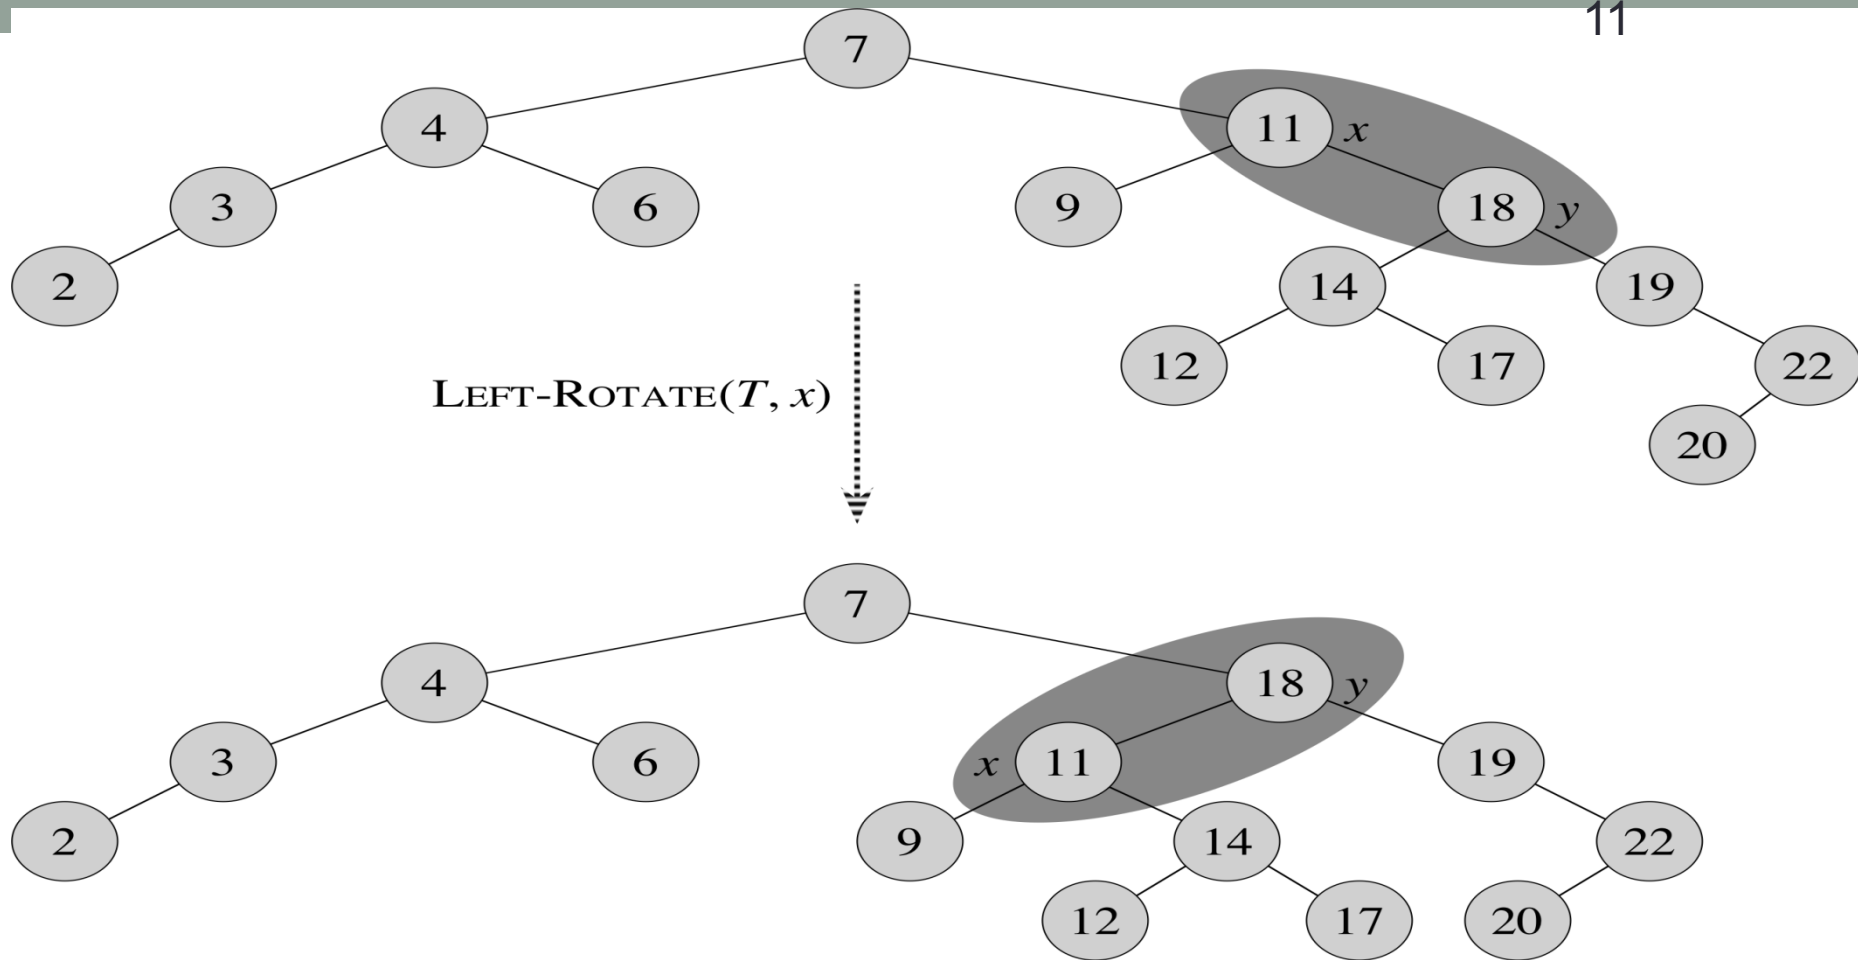
- Maintaining Red-Black tree properties upon insertion involves **rotations**



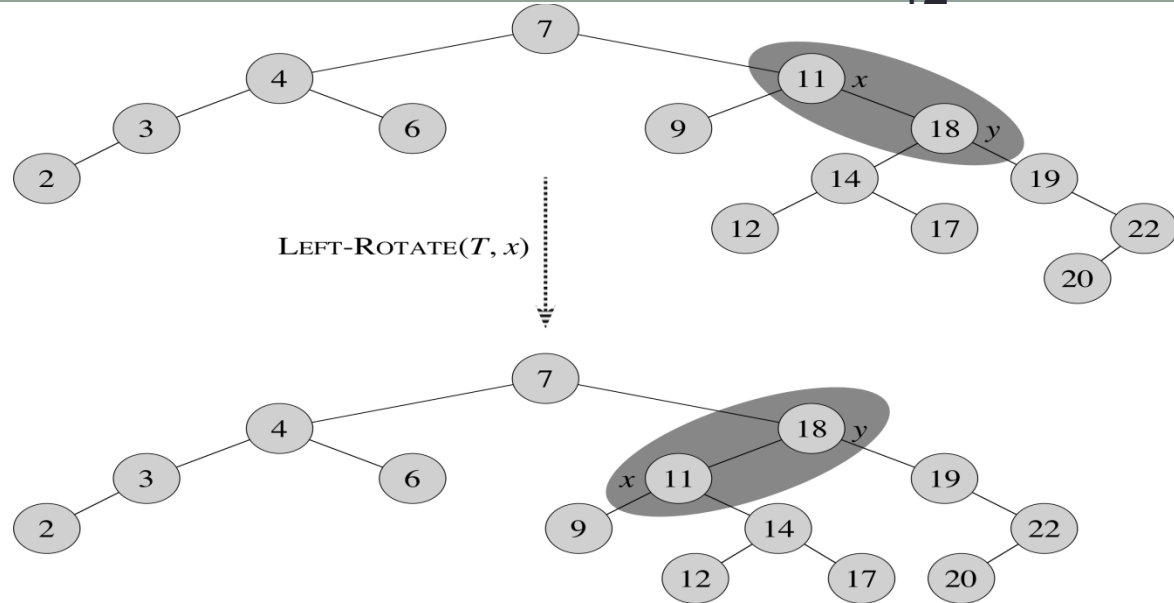
Left-Rotate:

- right child is not T.NIL.
- y is new root.
- x is y 's left child.
- y 's left child is x 's right child.

Rotation maintains inorder ordering of keys.
What happens to the relative subtree heights?



Rotations



- Before rotation: keys of x 's left subtree $\leq 11 \leq$ keys of y 's left subtree $\leq 18 \leq$ keys of y 's right subtree.
- Rotation makes y 's left subtree into x 's right subtree.
- After rotation: keys of x 's left subtree $\leq 11 \leq$ keys of x 's right subtree $\leq 18 \leq$ keys of y 's right subtree.
- Rotation is done by changing pointers. Time complexity = ??

Left-Rotate

LEFT-ROTATE(T, x)

```
1   $y = x.right$            // set y
2   $x.right = y.left$        // turn y's left subtree into x's right subtree
3  if  $y.left \neq T.nil$ 
4       $y.left.p = x$ 
5   $y.p = x.p$              // link x's parent to y
6  if  $x.p == T.nil$ 
7       $T.root = y$ 
8  elseif  $x == x.p.left$ 
9       $x.p.left = y$ 
10 else  $x.p.right = y$ 
11  $y.left = x$            // put x on y's left
12  $x.p = y$ 
```

Time complexity?

Insertion and deletion

- SEARCH, MIN, MAX, SUCCESSOR, PREDECESSOR are the same as with binary search tree.
- If we **insert**, what color to make the new node?
 - Red? Might violate property 4. (if node is red, both children are black)
 - Black? Might violate property 5 (all paths to descendants have same number of black nodes)
- If we **delete**, what color was the node that was removed?
 - Red? OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row.
 - Black? Could cause there to be two reds in a row (violating property 4), and can also cause a violation of property 5. Could also cause a violation of property 2 (root is black), if the removed node was the root and its child—which becomes the new root—was red.


R-B Insertion

- First, insert the new value as a **red** leaf into the R-B tree just like a BST.
 - $\text{RB-Insert}(T, z)$
- This may cause a **red-red** violation or a **red** root violation, thus needing to be fixed.
 - $\text{RB-Insert-Fixup}(T, z)$

THE LITTLE PICTURE




The idea of fixup

- You start with one of these



You want this




This is CASE 3



Rotate these




The idea of fixup

- You start with one of these:
- Remove the Zig-Zag
- Then rotate
- Call this Case 2



The idea of fixup (call this case 1)


- You start with this



Look for the red uncle!

You can change to this:
It does not change the black height.
The next node can be added without further adjustment
You might still have a red-red violation (you have moved the problem up the tree).

Say 10 is inserted



What case is this?

What is the resulting tree
After the fix-up procedure?

RB-INSERT(*T*, *z*)

```
1  y = T.nil
2  x = T.root
3  while x ≠ T.nil
4      y = x
5      if z.key < x.key
6          x = x.left
7      else x = x.right
8  z.p = y
9  if y == T.nil
10     T.root = z
11  elseif z.key < y.key
12     y.left = z
13  else y.right = z
14  z.left = T.nil
15  z.right = T.nil
16  z.color = RED
17  RB-INSERT-FIXUP(T, z)
```

Lines 3-7 search for the correct leaf position.

Note that **nil** is not the same as NULL.

Remember that the newly inserted node *z* is first colored red.

What are 4 differences between Tree-Insert and RB-Insert?

Unfortunately, the Fixup is more complicated!

RB-INSERT-FIXUP(T, z)

```

1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$ 
6               $y.color = BLACK$ 
7               $z.p.p.color = RED$ 
8               $z = z.p.p$ 
9          else if  $z == z.p.right$ 
10              $z = z.p$ 
11             LEFT-ROTATE( $T, z$ )
12              $z.p.color = BLACK$ 
13              $z.p.p.color = RED$ 
14             RIGHT-ROTATE( $T, z.p.p$ )
15         else (same as then clause
              with “right” and “left” exchanged)
16   $T.root.color = BLACK$ 

```

Note that y is
the **UNCLE** of z .

if red uncle

color parent black // case 1

color uncle black // case 1

color grandparent red // case 1

move problem to gp // case 1

if right child with black uncle

move problem up to parent // case 2

rotate on parent node // case 2

z is a left child with black uncle // case 3

change parent & gp colors // case 3

rotate right // case 3

RB-INSERT-FIXUP(T, z)

```

1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9          else if  $z == z.p.right$ 
10              $z = z.p$ 
11             LEFT-ROTATE( $T, z$ )
12              $z.p.color = \text{BLACK}$ 
13              $z.p.p.color = \text{RED}$ 
14             RIGHT-ROTATE( $T, z.p.p$ )
15         else (same as then clause
                with “right” and “left” exchanged)
16   $T.root.color = \text{BLACK}$ 

```

```


if (uncle.color == red)
{
    # Handle case
}
else
{
    if (z == z.p.right)
    {
        # Handle case 2
    }
    # Handle case 3
}

```


Three Cases, summary

- Case 1: Red uncle. Recolor parent and grandparent. Moves problem to grandparent
- Case 2: Black uncle and zig-zagging. Move problem to parent. Rotate with new parent to remove zig-zag.
- Case 3: Black uncle without zig-zag. Recolor parent and grandparent. Rotate on parent-grandparent.


(a)




(b)



(c)



(d)




Case 1: If **red uncle**,
Re-color above generations.
Also, move problem to gp.
Note that y is z's uncle


Case 2: Black uncle and
Zig-Zagging
(z is a right child
and z's parent is a left child):
Move problem to parent.
Rotate on parent to remove
zig-zag

Case 3: Black uncle
Without zig-zag
(z is a left child
and z's parent is a left child):
Recolor parent to black and gp to red.
Rotate on parent-gp to balance
Terminal Case!


Another Example: Insert 15




Case 1




Case 2




Case 3



Insert 50



Insert 65



Interview Questions

- Why is a new node set to red and not black?
- Show the red-black trees that result after successively inserting the keys 41, 38, 31, 12, 19, 8 into an initially empty red-black tree.

41, 38, 31, 12, 19, 8